

Bitwise Tips :

Left shift ($<<$):

- shifting left is equivalent to multiplication by power of 2, for example:

$$\bullet 6 << 1 = 6 \times 2$$

$$\bullet 6 << 3 = 6 \times (2^3) = 6 \times 8 = 48$$

Right shift ($>>$):

- shifting right is equivalent to division by power of 2, for example:

$$\bullet 12 >> 1 = 12 : (2^1) = 6$$

$$\bullet 12 >> 2 = 12 : (2^2) = 3$$

$$x = \sum_{k=0}^{w-1} x_k 2^k \quad (\text{unsigned integer})$$

$$x = \left(\sum_{k=0}^{w-2} x_k 2^k \right) - \boxed{x_{w-1} 2^{w-1}} \quad \begin{matrix} \text{sign bit} \\ \text{sign integer} \end{matrix}$$

- Unsigned integer, example:

8 bit word $0b10010110$ represents the unsigned value, $150 = 2 + 4 + 16 + 128$

- Signed integer (two's complement), example:

8 bit word $0b10010110$ represents the signed value $-106 = 2 + 4 + 16 - 128$

- We have $0b00\dots0 = 0$

- What is the value of $x = 0b11\dots1$?

$$x = \left(\sum_{k=0}^{w-2} x_k 2^k \right) - x_{w-1} 2^{w-1}$$

$$= \left(\sum_{k=0}^{w-2} 2^k \right) - 2^{w-1}$$

$$= (2^{w-1} - 1) - 2^{w-1}$$

$$= -1$$

Complementary Relationship

$$x + \sim x = -1 \Rightarrow -x = \sim x + 1$$

Exp:

$$x = \text{ob}011011000$$

$$\sim x = \text{ob}100100111$$

$$-x = \text{ob}100101000$$

* The prefix "Ob" designates a Boolean constant.

To translate from hex to binary, translate each hex digit to its binary equivalent, and concatenate the bits.

Example: 0x1DEC1DE2CODE4FOOD is
→ the prefix 0x designates a hex constant

1101 1110 1100 0001 1101 1110 0010 1100 0000 1101 1110 0100
↓ ↓ ↓ ↓ ↓ ↓ ↓ ↓ ↓ ↓ ↓ ↓ ↓
D E C 1 D E 2 C O D E 4

1111 0000 0000 1101
↓ ↓ ↓ ↓
F 0 0 D

Bitwise operators :

Operator	Description
&	AND
	OR
\wedge	XOR (exclusive OR)
~	NOT (one's complement)
\ll	Shift left
\gg	Shift right

- Examples (8-bit word)

$$A = \text{Ob } 1011\ 0011$$

$$B = \text{Ob } 0110\ 1001$$

$$A \& B = \text{Ob } 0010\ 0001$$

$$A | B = \text{Ob } 1111\ 1011$$

$$A ^ B = \text{Ob } 1101\ 1010$$

$$\sim A = \text{Ob } 0100\ 1100$$

$$A \gg 3 = \text{Ob } 0001\ 0110$$

$$A \ll 2 = \text{Ob } 1100\ 1100$$

PROBLEMS :

1. Set k th bit in a word $x + 1$.

- IDEA :

* SHIFT and OR.

$$y = x \mid (1 \ll k)$$

- EXAMPLE :

$$k = 7$$

x	$1 \ll k$	$x \mid (1 \ll k)$	7th
10111101	00000000	1011110111101101	11011101

2. Clear the k th bit in a word x .

- IDEA :

* SHIFT, complement, and AND

$$y = x' \& \sim(1 \ll k)$$

- EXAMPLE :

$$k = 7$$

→ 7th bit

x	$10111101\boxed{1}1101101$
$1 \ll k$	0000000010000000
$\sim(1 \ll k)$	1111111101111111
$x \& \sim(1 \ll k)$	1011110101101101

3. Toggle / Flip the k th bit in a word x

- IDEA:

* SHIFT and XOR

$$y = x \wedge (1 \ll k)$$

- EXAMPLE ($0 \rightarrow 1$):

$$k = 7$$

x	$10111101\boxed{1}1101101$
$1 \ll k$	0000000010000000
$x \wedge (1 \ll k)$	1011110111101101

4. Extract a bit field from a word x

- IDEA:

* mask and shift

$$(x \& \text{mask}) \gg \text{shift}$$

- EXAMPLE:
shift = 7

x
mask
 $x \& mask$
 $x \& mask \gg shift$

1011110101101101
0000011110000000
0000010100000000
0000000000001010

- this is a good trick to know if you're working with compressed or encoded data.

5. Set a bit field in a word x to a value y.

- ~~NEA~~:

* invert mask to clear, and OR the shifted value

$x = (x \& \sim \text{mask}) | (y \ll \text{shift})$

- EXAMPLE:

shift = 7 For safety's sake. $((y \ll \text{shift}) \& \text{mask})$

x
y
mask
 $x \& \sim \text{mask}$
 $(x \& \sim \text{mask}) | (y \ll \text{shift})$

1011110101101101
0000000000000001
0000111100000000
1011100001101101
101100111101101

6. Ordinary swap: swap two integers x and y .
- the standard way of doing this is to use a temporary variable:

$$t = x$$

$$x = y$$

$$y = t$$

* we can do this by using bit tricks,
swapping x and y without using a temporary.

$$x = x \wedge y \rightarrow \text{Mask with 1's where bits differ.}$$

$$y = x \wedge y \rightarrow \text{Flip the bits in } y \text{ that differ from } x.$$

$$x = x \wedge y \quad \downarrow \text{Flip the bits in } x \text{ that differ from } y.$$

For example:

x	10111101	10010011	10010011	00101110
y	00101110	00101110	10111101	10111101

- Why it works:

XOR is its own inverse:

$$(x \wedge y) \wedge y \Rightarrow x$$

x	y	x^y	$(x^y)^y$
0	0	0	0
0	1	1	0
1	0	1	1
1	1	0	1

* when we xor something twice, it just cancel out and we get back the original thing.

- Performance :

Poor at exploiting "instruction-level parallelism (ILP)". ("the naive approach is faster").

7. Find the minimum ' ϵ ' of two integers x and y .

The standard approach will be :

if ($x < y$):

$$\epsilon = x$$

else :

$$\epsilon = y$$

- Performance:

one performance problem with this code is that there is a branch in this code (modern machines will do branch prediction, and for whatever branch predicts the code to take it's going to do prefetching and execute some of the instructions in advance). But the problem is, if it mispredicts the branch, it does a lot of wasted work, and the processor has to empty the pipeline and undo all of the work that it did.

- Caveat:

the compiler is usually smart enough to optimize away the unpredictable branch, but maybe not.

? Is there a way to do a minimum without using a branch

$$r = y \wedge ((x \wedge y) \& -(x < y))$$

Why it works:

- The C language represents the Booleans TRUE and FALSE with the integers 1 and 0.
- If $x < y$, then $-(x < y) \Rightarrow -1$, which is all 1's in two's complement representation. Therefore, we have $y^1(x^1y) \Rightarrow x$
- If $x \geq y$, then $-(x < y) \Rightarrow 0$. Therefore, we have $y^0 \Rightarrow y$.

Info! Modern compilers can perform this optimization better than you can. (the branchless version is usually slower than the branching).

- * Why learn bit hacks if they don't even work?
 - Because the compiler does them, and it will help to understand what the compiler is doing when you look at the assembly code.
 - Because sometimes the compiler doesn't optimize, and you have to do it yourself by hand.
 - Because many bit hacks for words extend naturally to bit and word hacks for vectors.
 - Because these tricks arise in other domains.
 - Because they're fun! YESSS :)

8. Modular Addition : compute $(x+y) \bmod n$,
assuming that $0 \leq x < n$ and $0 \leq y < n$.

$$z = (x+y) \% n \quad (\text{division is expensive, unless by a power of 2})$$

$$\begin{aligned} z &= x + y && (\text{Unpredictable}) \\ z &= z \text{ if } z < n \text{ else } z - n && (\text{branch is expensive}) \end{aligned}$$

$$\begin{aligned} z &= x + y \\ z &= z - (n \& -(z \geq n)) \Rightarrow \text{same task as minimum} \end{aligned}$$

9. Round up to a Power of 2

Compute $2^{\lg n} \rightarrow$ Notation
 $\lg n = \log_2 n$

uint64_t n; Bit $\lceil \lg n \rceil - 1$ Example:
: flip the rightmost value of 1 with 0,
: and rest with 1
-- n; (decrement n)
n |= n >> 1

0010000001010000
0010000001001111
0011000001101111

$$n \mid = n \gg 2$$

$$n \mid = n \gg 4$$

$$n \mid = n \gg 8$$

$$n \mid = n \gg 16$$

$$n \mid = n \gg 32$$

Populate all
bits to the
right with 0

0011110001111111

0111111111111111

:

:

:

:

\uparrow we add one to that
 \hookrightarrow set bit $\lceil \log n \rceil$

0100000000000000

Why decrement?

- if n is already a power of 2 and if we don't decrement n , this isn't going to work because the $\log n - 1$ bit isn't set. But if we decrement n , then it's going to guarantee us that the $\log n - 1$ bit is set so that we can propagate that to the right.

10. Compute the mask of the least-significant 1 in word x .

- EXAMPLE: We want a mask that has a 1 in only the position of the least significant 1 in x , and 0's everywhere else.

$$r = x \& (-x)$$

- Exp:

X	00100000 1010000
-X	11011111 0110000
$x \& (-x)$	00000000 0010000

- Why it works

- the binary representation of $-x$ is $(nx) + 1$

- Question

How do you find the index of the bit, $\lg x$?

11. Log Base 2 of a Power of 2:

Compute $\lg x$, where x is a power of 2

const uint64_t deBruijn = 0x022fdd63cc95386d

const int convert[64] = {

k=8 0, 1, 2, 53, 3, 7, 54, 27,
4, 38, 41, 8, 34, 55, 48, 28,
62, 5, 39, 46, 44, 42, 22, 9,
24, 35, 59, 56, 49, 18, 29, 11,

63, 52, 6, 26, 37, 40, 33, 47,
 61, 45, 43, 21, 23, 58, 17, 10,
 51, 25, 36, 32, 60, 20, 57, 16,
 50, 31, 19, 15, 30, 14, 13, 12
 ↴
 ↴

$\gamma = \text{Convert}[(x^* \text{ deBruijn}) \gg 58]$

- Why it works

A de Bruijn sequence s of length 2^k is a cyclic 0-1 sequence such that each of the 2^{k-1} strings of length k occurs exactly once as a substring of s.

- Example : $K=3$

All of the 8 possible substrings of length 3

0 0 0 1 1 1 0 1

0 0 0 ↙ Start with all 0's
 1 0 0 1

2 0 1 1

3 1 1 1

4 1 1 0

5 1 0 1

6 0 1 0

7

100

const int convert [8] = {0, 1, 6, 2, 7, 5, 4, 3}

For example:

this is the same as left shifting

$0b00011101 * 2^4 \Rightarrow 0b11010000$

$0b11010000 \gg 5 \Rightarrow 6$

convert [6] $\Rightarrow 4$

110 appears at the beginning of the sequence

we want to extract this out, and we
right shifting five position $\Rightarrow 110 = 6$

Performance

Limited by multiply and table look up.

12. Population Count I:

Count the number of 1 bits in a word x.

for ($r=0$; $x \neq 0$; $++r$) Repeatedly

$x \&= x - 1$

eliminate the least-significant 1.

Example:

x	0010110111010000
$x - 1$	0010110111001111
$x \& (x - 1)$	0010110111000000

Issue:

- fast if the population count is small,
summing time is proportional to the number
of bits in the word.

13. Population Count II:

```
static const int count[256] = 
{ 0, 1, 1, 2, 1, 2, 2, 3, 1, ..., 8 }
for (int k=0; x != 0; x >>= 8)
    k += count[x & 0xFF];
    K times
     $x = \underbrace{xx \dots x}_{k}$ 
```

14. Population Count III:

// Create masks

$$\begin{aligned} M_5 &= \sim((-1) \ll 32); \quad // 0^{32, 32} \\ M_4 &= M_5 \wedge (M_5 \ll 16); // (0^{16, 16})^2 \\ M_3 &= M_4 \wedge (M_4 \ll 8); // (0^8, 8)^4 \\ M_2 &= M_3 \wedge (M_3 \ll 4); // (0^4, 4)^8 \\ M_1 &= M_2 \wedge (M_2 \ll 2); // (0^2, 2)^16 \end{aligned}$$

$$M0 = M1 \wedge (M1 \ll 1); // (0^1)^{32}$$

// Compute pop count

Avoid overflow

$$x = ((x \gg 1) \& M0) + (x \& M0)$$
$$x = (((x \gg 2) \& M1) + (x \& M1))$$
$$x = (((x \gg 4) + x) \& M2)$$
$$x = (((x \gg 8) + x) \& M3)$$
$$x = (((x \gg 16) + x) \& M4)$$
$$x = (((x \gg 32) + x) \& M5)$$

No worry about overflow

Performance: $O(\lg w)$ time, where $w = \text{word length}$